


Chapter 9: Virtual-Memory Management






Chapter 9: Virtual-Memory Management

- n Background
- n Demand Paging
- n Copy-on-Write
- n Page Replacement
- n Allocation of Frames
- n Thrashing
- n Memory-Mapped Files
- n Allocating Kernel Memory
- n Other Considerations
- n Operating-System Examples






Objectives

- n To describe the benefits of a virtual memory system
- n To explain the concepts of demand paging, page-replacement algorithms, and allocation of page frames
- n To discuss the principle of the working-set model






Background

- n Code needs to be in memory to execute, but entire program rarely used
 - | Error code, unusual routines, large data structures
- n Entire program code not needed at same time
- n Consider ability to execute partially-loaded program
 - | Program no longer constrained by limits of physical memory
 - | Program and programs could be larger than physical memory






Background

- n **Virtual memory** – separation of user logical memory from physical memory
 - | Only part of the **program needs to be in memory for execution**
 - | Logical address space can therefore be much larger than physical address space
 - | Allows address spaces to be **shared by several processes**
 - | Allows for more efficient process creation
 - | More programs running **concurrently**
 - | Less I/O needed to load or swap processes
- n Virtual memory can be implemented via:
 - | Demand paging
 - | Demand segmentation







Virtual Memory That is Larger Than Physical Memory





Virtual-address Space






Virtual Address Space


- n Enables **sparse** address spaces with holes left for growth, dynamically linked libraries, etc
- n System libraries shared via mapping into virtual address space
- n Shared memory by mapping pages read-write into virtual address space
- n Pages can be shared during `fork()`, speeding process creation





Shared Library Using Virtual Memory






Demand Paging


- n Could bring entire process into memory at load time
- n Or bring a page into memory only when it is needed
 - | Less I/O needed, no unnecessary I/O
 - | Less memory needed
 - | Faster response
 - | More users
- n Page is needed \Rightarrow reference to it
 - | invalid reference \Rightarrow abort
 - | not-in-memory \Rightarrow bring to memory
- n **Lazy swapper** – never swaps a page into memory unless page will be needed
 - | Swapper that deals with pages is a **pager**





Transfer of a Paged Memory to Contiguous Disk Space





Valid-Invalid Bit


- n With each page table entry a valid–invalid bit is associated (**v** \Rightarrow in-memory – **memory resident**, **i** \Rightarrow not-in-memory)
- n Initially valid–invalid bit is set to **i** on all entries
- n Example of a page table snapshot:

Frame #	valid-invalid bit
	v
	v
	v
	v
	i
....	
	i
	i

page table

- n During address translation, if valid–invalid bit in page table entry is **i** \Rightarrow page fault





Page Table When Some Pages Are Not in Main Memory

0	A
1	B
2	C
3	D
4	E
5	F
6	G
7	H


logical
memory


valid-invalid bit		
frame		
0	4	v
1		i
2	6	v
3		i
4		i
5	9	v
6		i
7		i

page table

0	
1	
2	
3	
4	A
5	
6	C
7	
8	
9	F
10	
11	
12	
13	
14	
15	

physical memory





Page Fault

- n If there is a reference to a page, first reference to that page will trap to operating system:

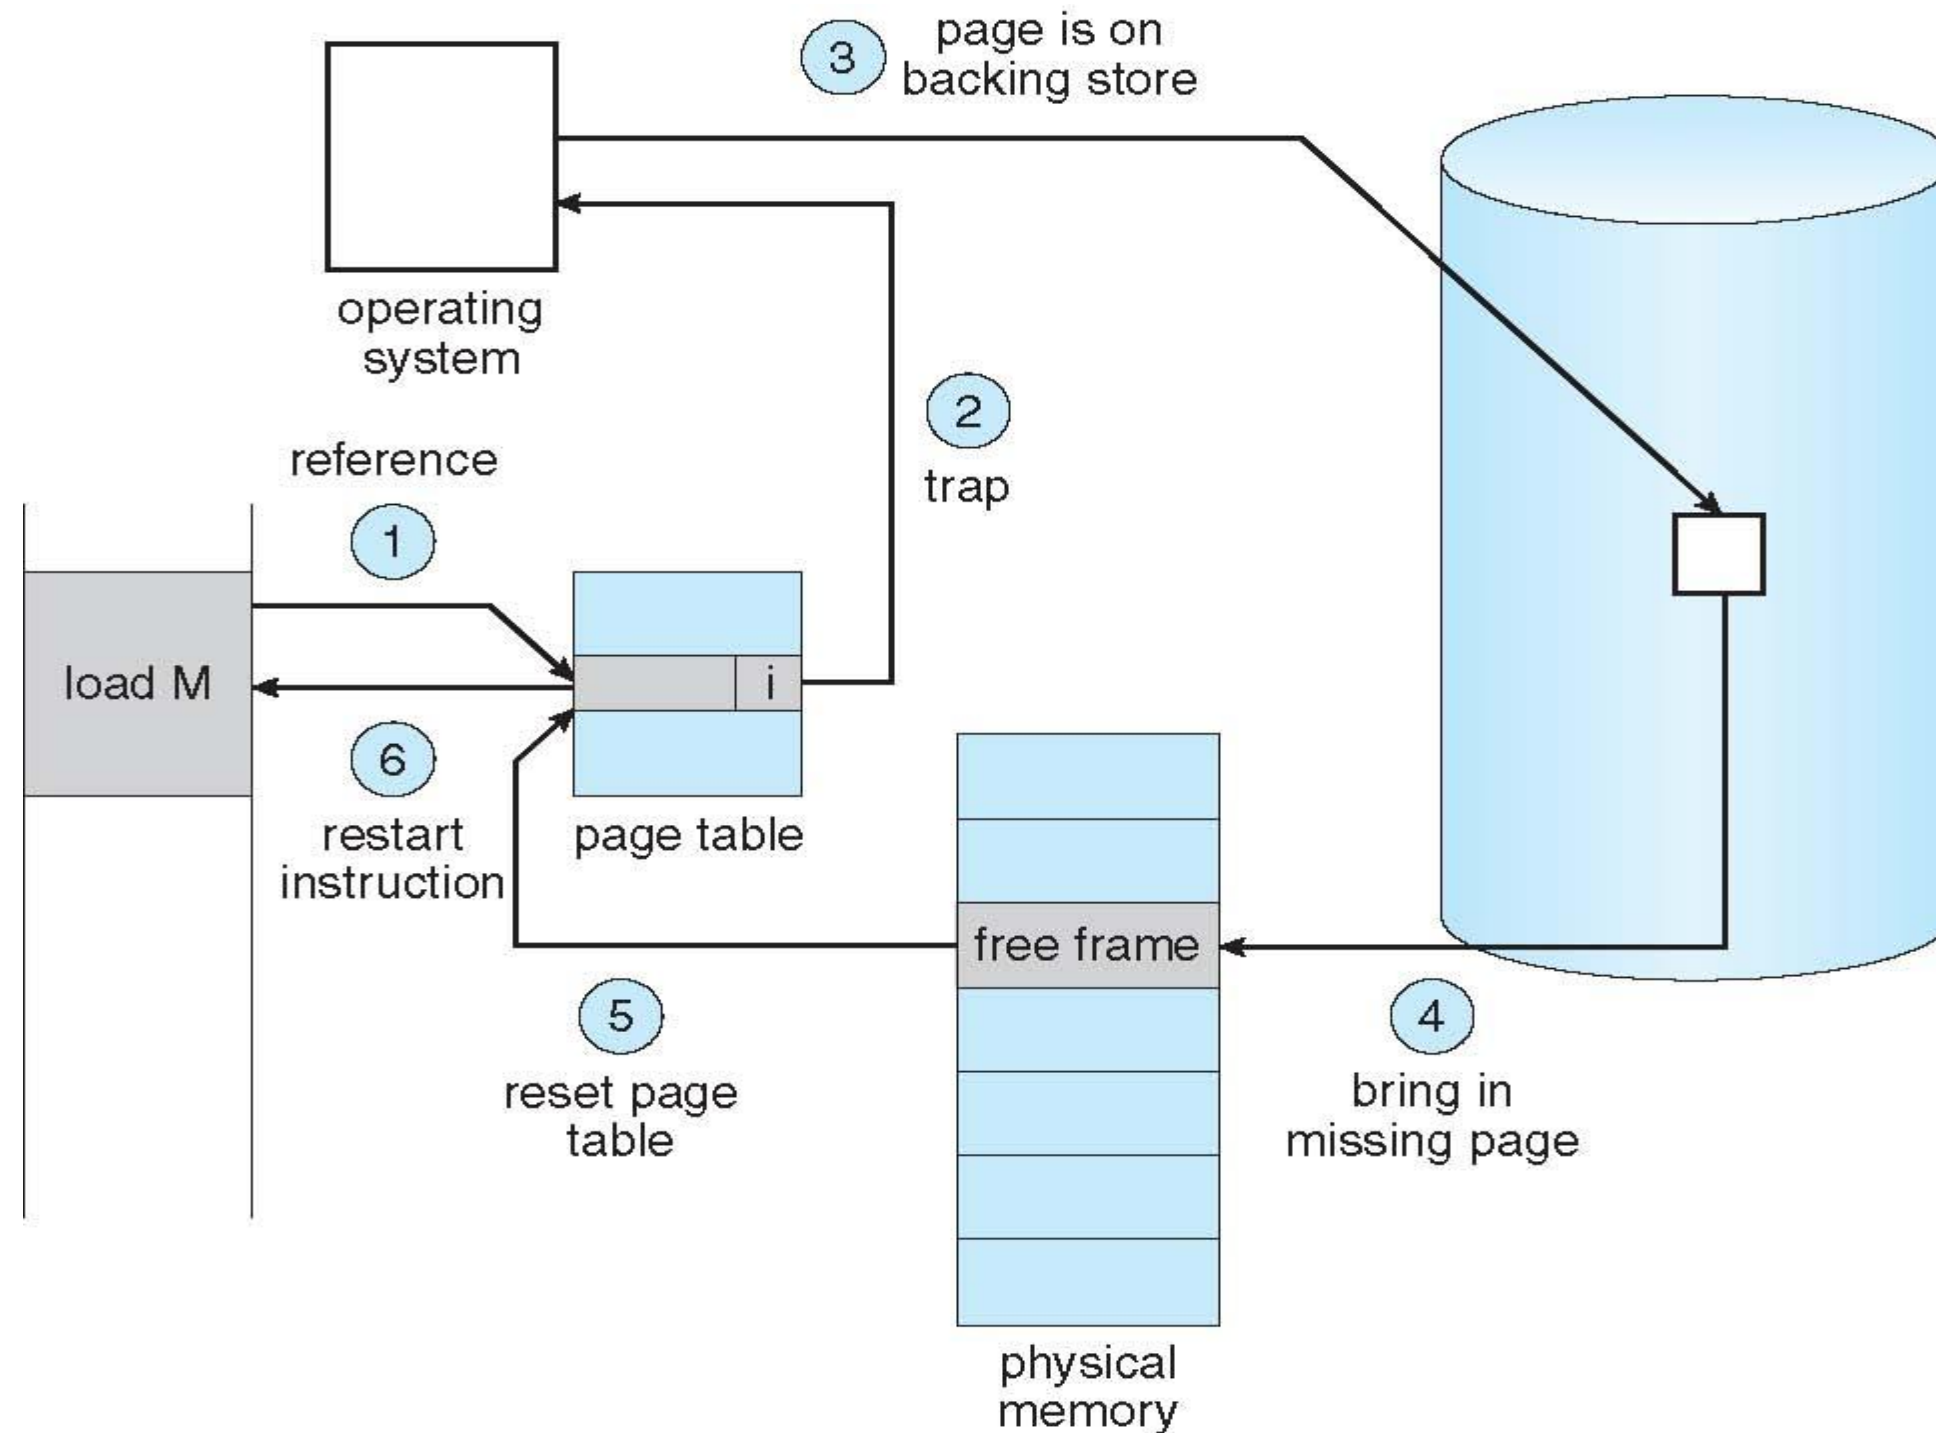
page fault


1. Operating system looks at another table to decide:
 - | Invalid reference \Rightarrow abort
 - | Just not in memory
2. Get empty frame
3. Swap page into frame via scheduled disk operation
4. Reset tables to indicate page now in memory
Set validation bit = **v**
5. Restart the instruction that caused the page fault





Steps in Handling a Page Fault






Aspects of Demand Paging

- n Extreme case – start process with no pages in memory
 - | OS sets instruction pointer to first instruction of process, non-memory-resident -> page fault
 - | And for every other process pages on first access
 - | Pure demand paging
- n Actually, a given instruction could access multiple pages -> multiple page faults
 - | Pain decreased because of locality of reference
- n Hardware support needed for demand paging
 - | Page table with valid / invalid bit
 - | Secondary memory (swap device with swap space)
 - | Instruction restart






Instruction Restart

- n Consider an instruction that could access several different locations
 - | For example, consider the IBM System 360/370 MVC (move character) instruction, which can move up to 256 bytes from one location to another (possibly overlapping) location.
 - | If either block (source or destination) straddles a page boundary, a page fault might occur after the move is partially done. In addition, if the source and destination blocks overlap, the source block may have been modified, in which case we cannot simply restart the instruction.
- n This problem can be solved in two different ways.
 - | In one solution, the microcode computes and attempts to access both ends of both blocks. If a page fault is going to occur, it will happen at this step, before anything is modified. The move can then take place; we know that no page fault can occur, since all the relevant pages are in memory.
 - | The other solution uses temporary registers to hold the values of overwritten locations. If there is a page fault, all the old values are written back into memory before the trap occurs. This action restores memory to its state before the instruction was started, so that the instruction can be repeated.






Performance of Demand Paging

- n Stages in Demand Paging
 1. Trap to the operating system
 2. Save the user registers and process state
 3. Determine that the interrupt was a page fault
 4. Check that the page reference was legal and determine the location of the page on the disk
 5. Issue a read from the disk to a free frame:
 1. Wait in a queue for this device until the read request is serviced
 2. Wait for the device seek and/or latency time
 3. Begin the transfer of the page to a free frame
 6. While waiting, allocate the CPU to some other user
 7. Receive an interrupt from the disk I/O subsystem (I/O completed)
 8. Save the registers and process state for the other user
 9. Determine that the interrupt was from the disk
 10. Correct the page table and other tables to show page is now in memory
 11. Wait for the CPU to be allocated to this process again
 12. Restore the user registers, process state, and new page table, and then resume the interrupted instruction





Performance of Demand Paging (Cont.)

n Page Fault Rate $0 \leq p \leq 1$

| if $p = 0$ no page faults

| if $p = 1$, every reference is a fault

n Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} = & (1 - p) \times \text{memory access} \\ & + p (\text{page fault overhead} \\ & \quad + \text{swap page out} \\ & \quad + \text{swap page in} \\ & \quad + \text{restart overhead} \\ &) \end{aligned}$$





Demand Paging Example

- n Memory access time = 200 nanoseconds
- n Average page-fault service time = 8 milliseconds
- n
$$\begin{aligned} \text{EAT} &= (1 - p) \times 200 + p (8 \text{ milliseconds}) \\ &= (1 - p) \times 200 + p \times 8,000,000 \\ &= 200 + p \times 7,999,800 \end{aligned}$$
- n If one access out of 1,000 causes a page fault, then
EAT = 8.2 microseconds.
This is a slowdown by a factor of 40!!
- n If want performance degradation < 10 percent
 - | $220 > 200 + 7,999,800 \times p$
 $20 > 7,999,800 \times p$
 - | $p < .0000025$
 - | < one page fault in every 400,000 memory accesses






Demand Paging Optimizations

- n Copy entire process image to swap space at process load time
 - | Then page in and out of swap space
 - | Used in older BSD Unix

- n Demand page in from program binary on disk, but discard rather than paging out when freeing frame
 - | Used in Solaris and current BSD





Copy-on-Write

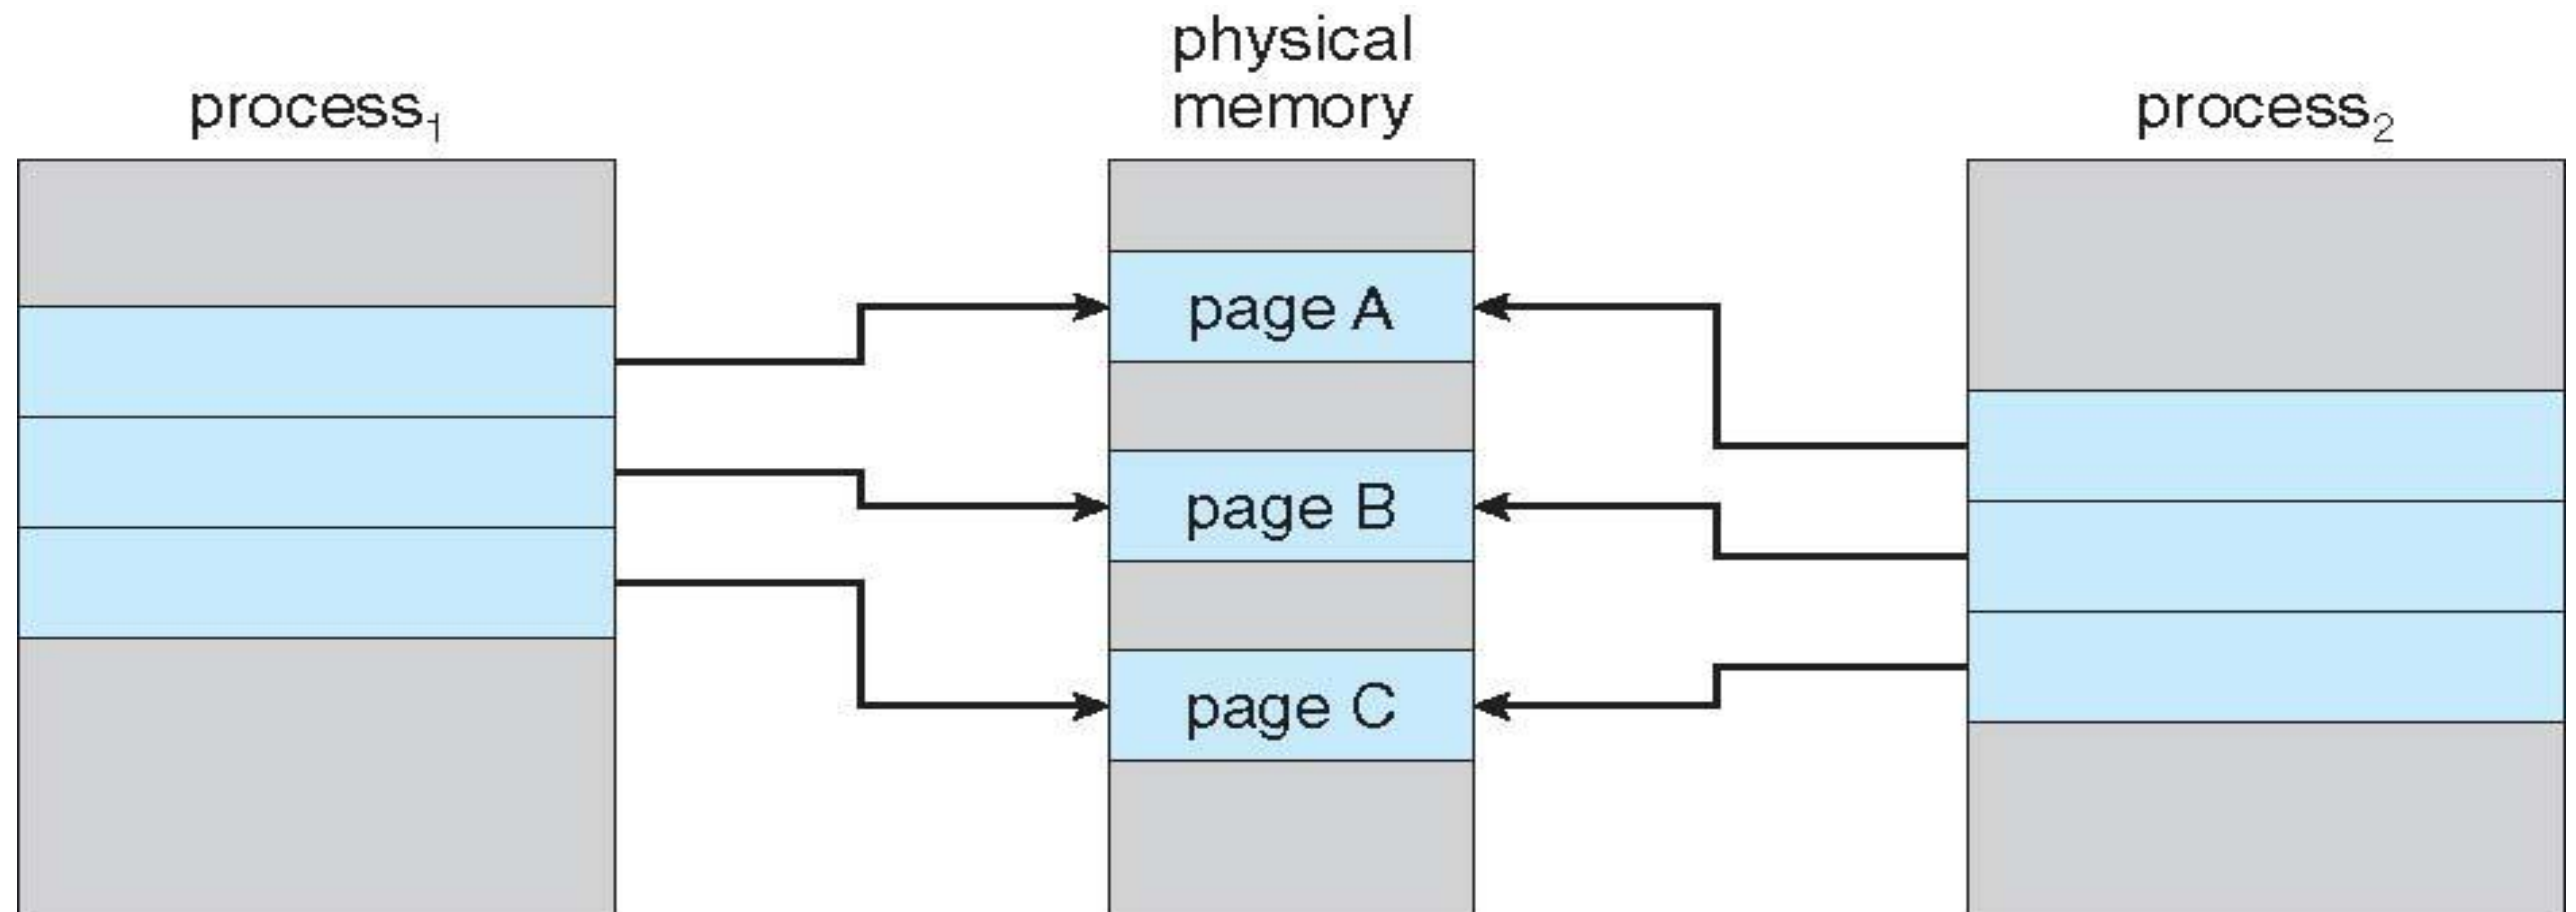
写入时复制


- n **Copy-on-Write** (COW) allows both parent and child processes to initially share the same pages in memory
 - | If either process modifies a shared page, only then is the page copied
- n COW allows more efficient process creation as only modified pages are copied
- n In general, free pages are allocated from a **pool** of **zero-fill-on-demand** pages
 - | Why zero-out a page before allocating it?
- n **vfork() variation on fork()** system call has parent suspend and child using copy-on-write address space of parent
 - | Designed to have child call exec()
 - | Very efficient





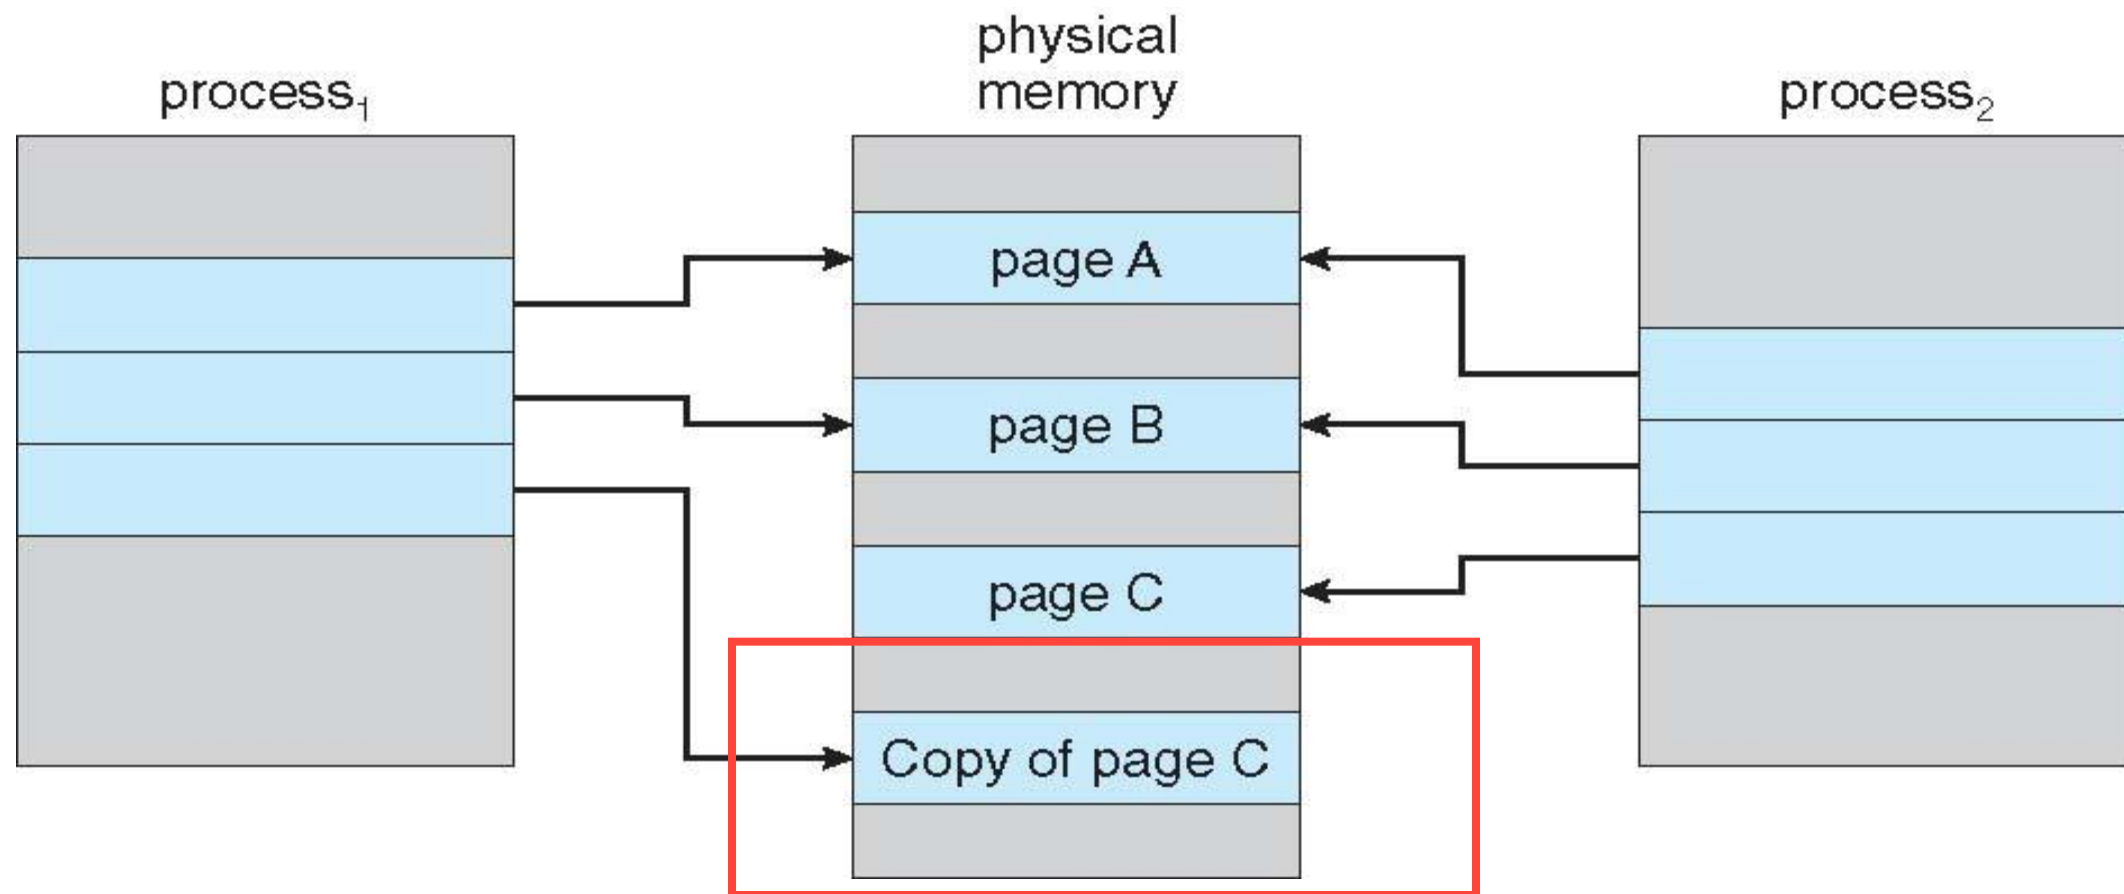
Before Process 1 Modifies Page C






After Process 1 Modifies Page C

如果有多个调用者 (Callers) 同时访问相同的资源 (如内存或者是磁盘上的数据存储)，他们会共同获取相同的指针指向相同的资源，直到某个调用者修改资源内容时，系统才会真正复制一份专用副本 (private copy) 给该调用者，而其他调用者所见到的最初的资源仍然保持不变。这过程对其他的调用者都是透明的 (transparently)。此做法主要的优点是如果调用者没有修改资源，就不会有副本 (private copy) 被创建，因此多个调用者只是读取操作时可以共享同一份资源。



读取时共享






What Happens if There is no Free Frame?

- n Used up by process pages
- n Also in demand from the kernel, I/O buffers, etc
- n How much to allocate to each?

- n **Page replacement** – find some page in memory, but not really in use, page it out
 - | Algorithm – terminate? swap out? replace the page?
 - | Performance – want an algorithm which will result in minimum number of page faults

- n Same page may be brought into memory several times








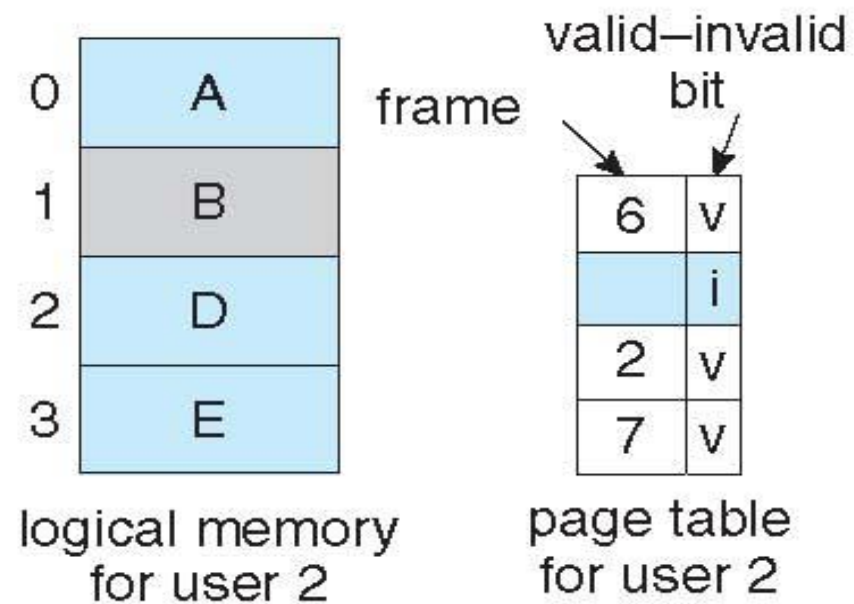

Page Replacement

- n Prevent **over-allocation** of memory by modifying page-fault service routine to include page replacement
- n Use **modify (dirty) bit** to **reduce overhead of page transfers – only modified pages are written to disk**
- n Page replacement completes separation between logical memory and physical memory – large virtual memory can be provided on a smaller physical memory





Need For Page Replacement






Basic Page Replacement

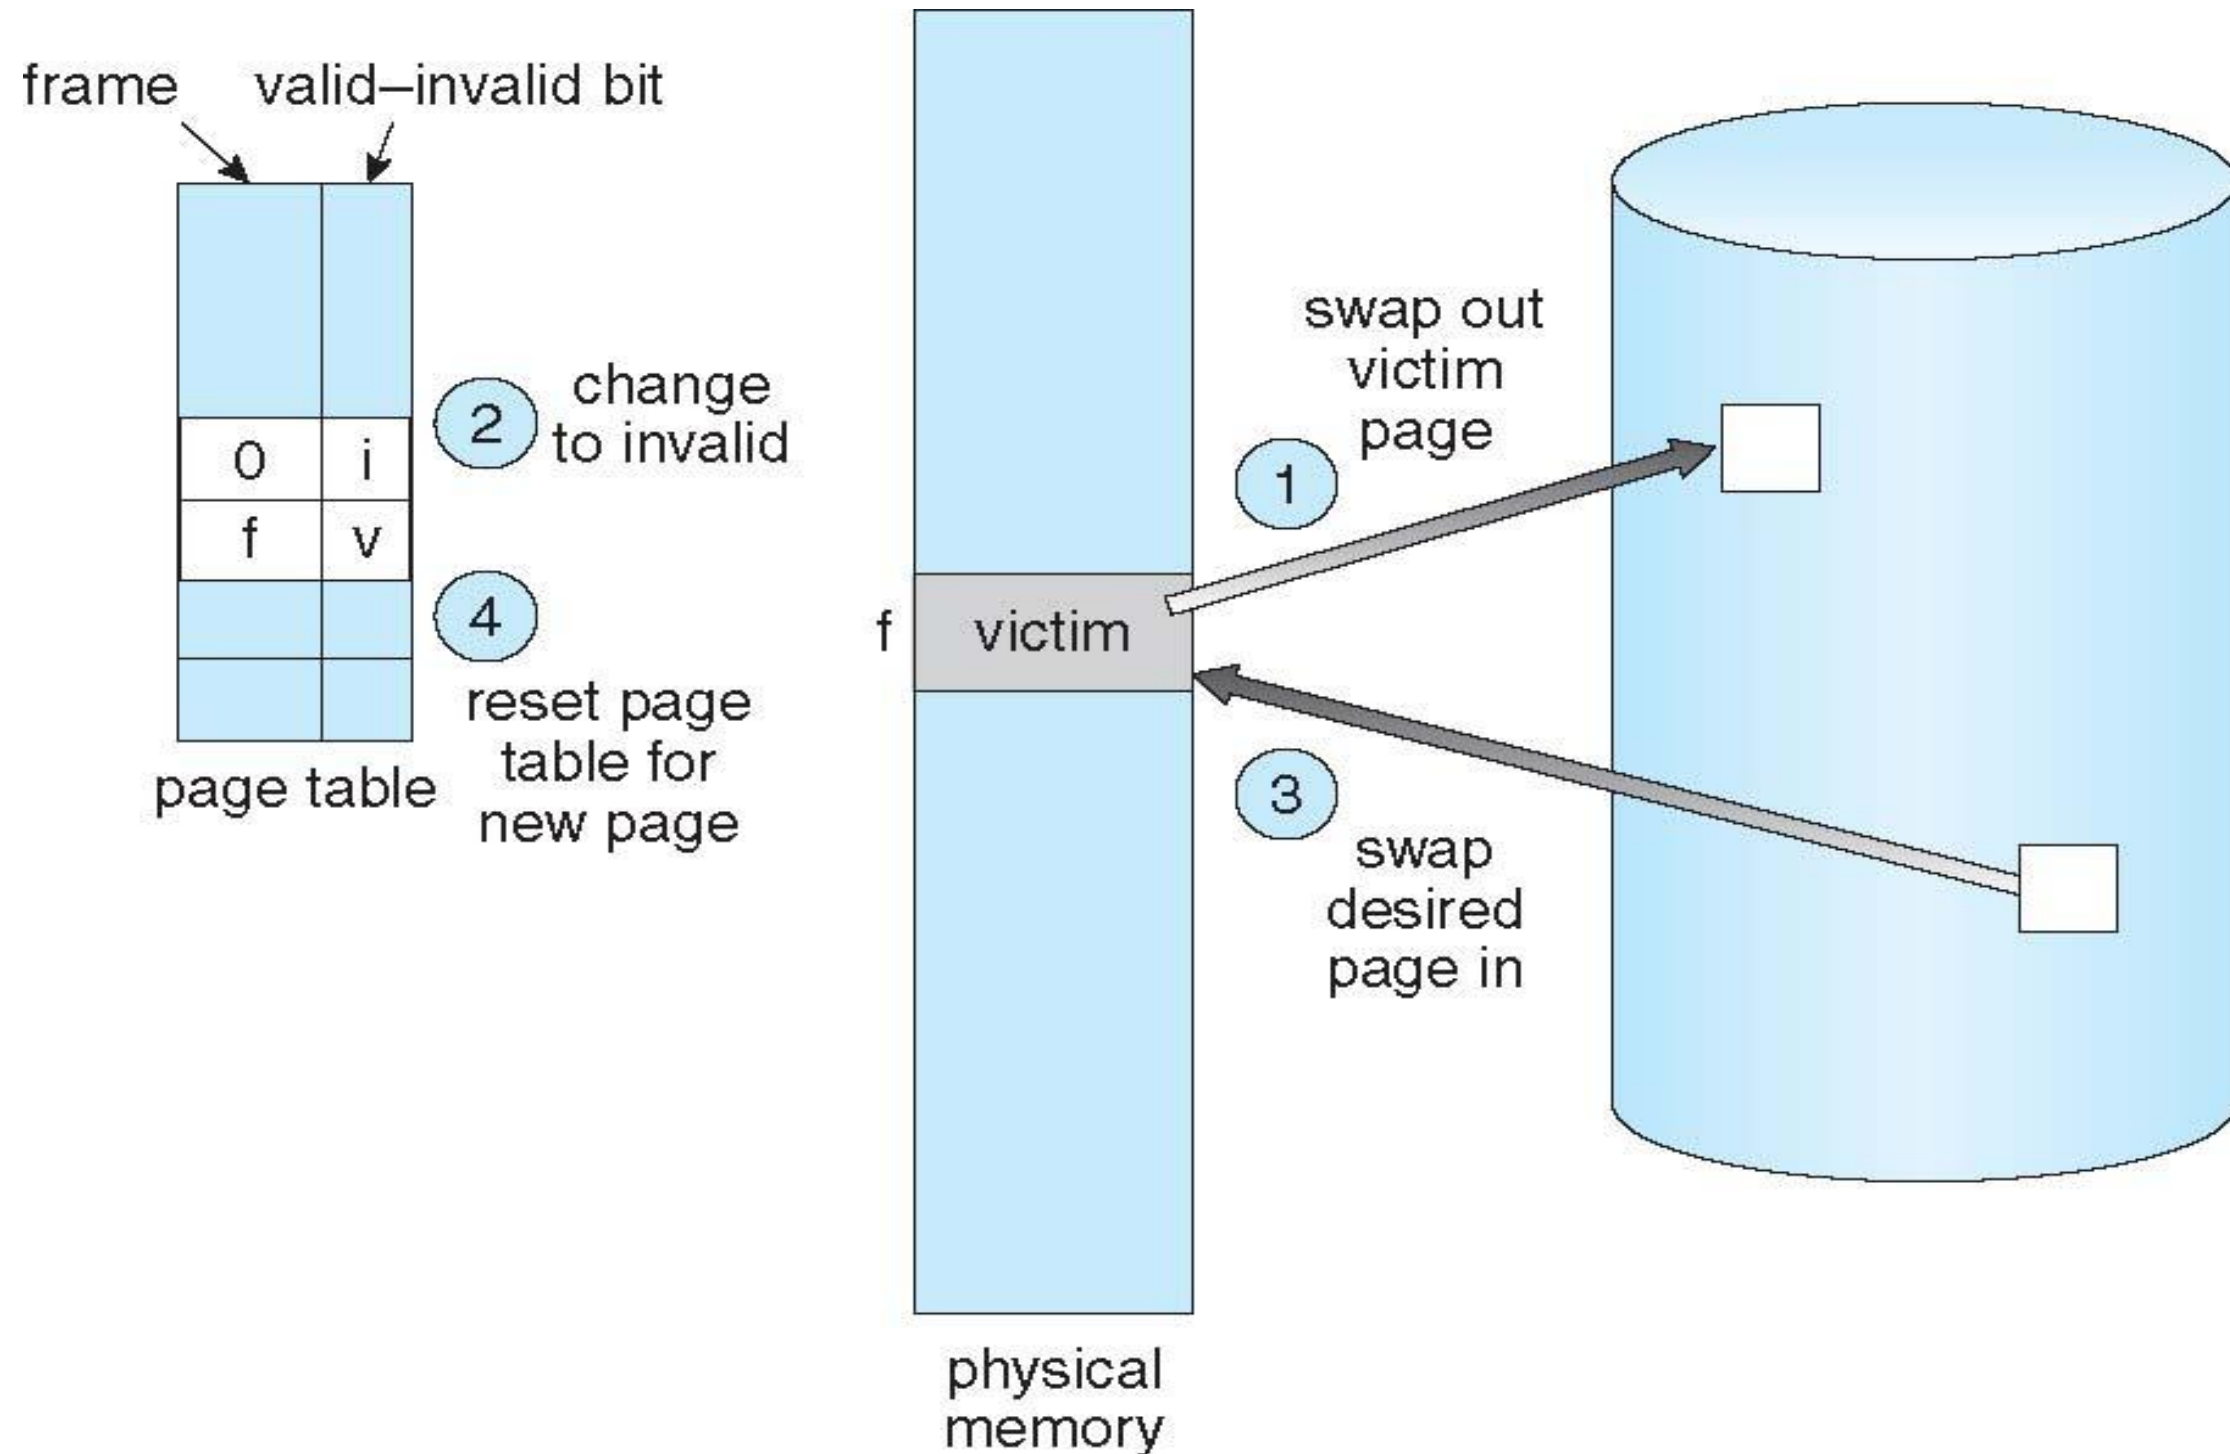
1. Find the location of the desired page on disk
2. Find a free frame:
 - If there is a free frame, use it
 - If there is no free frame, use a page replacement algorithm to select a **victim frame**
 - Write victim frame to disk if dirty
3. Bring the desired page into the (newly) free frame; update the page and frame tables
4. Continue the process by restarting the instruction that caused the trap

Note now potentially 2 page transfers for page fault – increasing EAT





Page Replacement





Page and Frame Replacement Algorithms


- n **Frame-allocation algorithm** determines
 - | How many frames to give each process
 - | Which frames to replace
- n **Page-replacement algorithm**
 - | Want lowest page-fault rate on both first access and re-access
- n Evaluate algorithm by running it on a particular string of memory references (reference string) and computing the number of page faults on that string
 - | String is just page numbers, not full addresses
 - | Repeated access to the same page does not cause a page fault
- n In all our examples, the reference string is

7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1





Graph of Page Faults Versus The Number of Frames





First-In-First-Out (FIFO) Algorithm

- n Reference string: 7,0,1,2,0,3,0,4,2,3,0,3,0,3,2,1,2,0,1,7,0,1
- n 3 frames (3 pages can be in memory at a time per process)

1	7	2	4	0	7
2	0	3	2	1	0
3	1	0	3	2	1

15 page faults

- n Can vary by reference string: consider 1,2,3,4,1,2,5,1,2,3,4,5
 - | Adding more frames can cause more page faults!
 - ▶ [Belady's Anomaly](#)

- n How to track ages of pages?
 - | Just use a FIFO queue





FIFO Page Replacement


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1


7	7	7	2																
	0	0	0																
		1	1																

page frames





FIFO Illustrating Belady's Anomaly






Optimal Algorithm

- n Replace page that will not be used for longest period of time
 - | 9 is optimal for the example on the next slide

- n How do you know this?
 - | Can't read the future

- n Used for measuring how well your algorithm performs






Optimal Page Replacement


reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1



page frames





Least Recently Used (LRU) Algorithm

- n Use past knowledge rather than future
- n Replace page that has not been used in the most amount of time
- n Associate time of last use with each page

reference string

7 0 1 2 0 3 0 4 2 3 0 3 2 1 2 0 1 7 0 1

7	7	7	2		2		4	4	4	0		1		1		1
	0	0	0		0		0	0	3	3		3		0		0
		1	1		3		3	2	2	2		2		2		7

page frames

- n 12 faults – better than FIFO but worse than OPT
- n Generally good algorithm and frequently used
- n But how to implement?





LRU Algorithm (Cont.)

n Counter implementation

- | Every page entry has a counter; every time page is referenced through this entry, copy the clock into the counter
- | When a page needs to be changed, look at the counters to find smallest value
 - ▶ Search through table needed

n Stack implementation

- | Keep a stack of page numbers in a double link form:
- | Page referenced:
 - ▶ move it to the top
 - ▶ requires 6 pointers to be changed
- | But each update more expensive
- | No search for replacement

n LRU and OPT are cases of **stack algorithms** that don't have Belady's Anomaly






Use Stack to Record The Most Recent Page References

reference string


4 7 0 7 1 0 1 2 1 2 7 1 2




stack
before
a



stack
after
b







LRU Approximation Algorithms

- n LRU needs special hardware and still slow
- n Reference bit
 - | With each page associate a bit, initially = 0
 - | When page is referenced bit set to 1
 - | Replace any with reference bit = 0 (if one exists)
 - ▶ We do not know the order, however
- n Second-chance algorithm
 - | Generally FIFO, plus hardware-provided reference bit
 - | Clock replacement
 - | If page to be replaced has
 - ▶ Reference bit = 0 -> replace it
 - ▶ reference bit = 1 then:
 - set reference bit 0, leave page in memory
 - replace next page, subject to same rules






Second-Chance (clock) Page-Replacement Algorithm



(a)



(b)





Counting Algorithms

- n Keep a counter of the number of references that have been made to each page
 - | Not common
- n **LFU Algorithm**: replaces page with smallest count
- n **MFU Algorithm**: based on the argument that the page with the smallest count was probably just brought in and has yet to be used





Page-Buffering Algorithms

- n Keep a pool of free frames, always
 - | Then frame available when needed, not found at fault time
 - | Read page into free frame and select victim to evict and add to free pool
 - | When convenient, evict victim
- n Possibly, keep list of modified pages
 - | When backing store otherwise idle, write pages there and set to non-dirty
- n Possibly, keep free frame contents intact and note what is in them
 - | If referenced again before reused, no need to load contents again from disk
 - | Generally useful to reduce penalty if wrong victim frame selected






Applications and Page Replacement

- n All of these algorithms have OS guessing about future page access
- n Some applications have better knowledge – i.e. databases
- n Memory intensive applications can cause double buffering
 - | OS keeps copy of page in memory as I/O buffer
 - | Application keeps page in memory for its own work
- n Operating system can given direct access to the disk, getting out of the way of the applications
 - | Raw disk mode
- n Bypasses buffering, locking, etc





Allocation of Frames

- n Each process needs minimum number of frames
- n Example: IBM 370 – 6 pages to handle SS MOVE instruction:
 - | instruction is 6 bytes, might span 2 pages
 - | 2 pages to handle from
 - | 2 pages to handle to
- n Maximum of course is total frames in the system
- n Two major allocation schemes
 - | fixed allocation
 - | priority allocation
- n Many variations





Fixed Allocation

- n Equal allocation – For example, if there are 100 frames (after allocating frames for the OS) and 5 processes, give each process 20 frames

- l Keep some as free frame buffer pool

- n Proportional allocation – Allocate according to the size of process

- l Dynamic as degree of multiprogramming, process sizes change

- s_i = size of process p_i

$$m = 64$$

- $S = \sum s_i$

$$s_1 = 10$$

- m = total number of frames


$$s_2 = 127$$

- a_i = allocation for $p_i = \frac{s_i}{S} \times m$

$$a_1 = \frac{10}{137} \times 64 \approx 5$$

$$a_2 = \frac{127}{137} \times 64 \approx 59$$






Priority Allocation

- n Use a proportional allocation scheme using priorities rather than size
- n If process P_i generates a page fault,
 - | select for replacement one of its frames
 - | select for replacement a frame from a process with lower priority number






Global vs. Local Allocation

- n **Global replacement** – process selects a replacement frame from the set of all frames; one process can take a frame from another
 - | But then process execution time can vary greatly
 - | But greater throughput so more common

- n **Local replacement** – each process selects from only its own set of allocated frames
 - | More consistent per-process performance
 - | But possibly underutilized memory






Non-Uniform Memory Access

- n So far all memory accessed equally
- n Many systems are NUMA – speed of access to memory varies
 - | Consider system boards containing CPUs and memory, interconnected over a system bus
- n Optimal performance comes from allocating memory “close to” the CPU on which the thread is scheduled
 - | And modifying the scheduler to schedule the thread on the same system board when possible
 - | Solved by Solaris by creating **lgroups**
 - ▶ Structure to track CPU / Memory low latency groups
 - ▶ Used my schedule and pager
 - ▶ When possible schedule all threads of a process and allocate all memory for that process within the lgroup







Thrashing

- n If a process does not have “enough” pages, the page-fault rate is very high
 - | Page fault to get page
 - | Replace existing frame
 - | But quickly need replaced frame back
 - | This leads to:
 - ▶ Low CPU utilization
 - ▶ Operating system thinking that it needs to increase the degree of multiprogramming
 - ▶ Another process added to the system
- n **Thrashing** \equiv a process is busy swapping pages in and out





Thrashing (Cont.)






Demand Paging and Thrashing


- n Why does demand paging work?
Locality model
 - | Process migrates from one locality to another
 - | Localities may overlap


- n Why does thrashing occur?
 Σ size of locality > total memory size
 - | Limit effects by using local or priority page replacement





Locality In A Memory-Reference Pattern





Working-Set Model

- n $\Delta \equiv$ working-set window \equiv a fixed number of page references
Example: 10,000 instructions
- n WSS_i (working set of Process P_i) =
total number of pages referenced in the most recent Δ (varies in time)
 - | if Δ too small will not encompass entire locality
 - | if Δ too large will encompass several localities
 - | if $\Delta = \infty \Rightarrow$ will encompass entire program
- n $D = \sum WSS_i \equiv$ total demand frames
 - | Approximation of locality
- n if $D > m \Rightarrow$ Thrashing, where m is available frames
- n Policy if $D > m$, then suspend or swap out one of the processes








Working-set model

page reference table

... 2 6 1 5 7 7 7 7 5 1 6 2 3 4 1 2 3 4 4 4 3 4 3 4 4 4 4 1 3 2 3 4 4 4 3 4 4 4 ...





Keeping Track of the Working Set


- n Approximate with interval timer + a reference bit
- n Example: $\Delta = 10,000$
 - | Timer interrupts after every 5000 time units
 - | Keep in memory 2 bits for each page
 - | Whenever a timer interrupts, copy and set the values of all reference bits to 0
 - | If one of the bits in memory = 1 \Rightarrow page in working set
- n Why is this not completely accurate?
- n Improvement = 10 bits and interrupt every 1000 time units






Page-Fault Frequency

- n More direct approach than WSS
- n Establish “acceptable” **page-fault frequency** rate and use local replacement policy
 - | If actual rate too low, process loses frame
 - | If actual rate too high, process gains frame





Working Sets and Page Fault Rates






Memory-Mapped Files

- n Memory-mapped file I/O allows file I/O to be **treated as routine memory access by mapping a disk block to a page in memory** use memory as a file to operate the disk.
- n A file is initially read using demand paging
 - | A page-sized portion of the file is read from the file system into a physical page
 - | Subsequent reads/writes to/from the file are treated as ordinary memory accesses
- n Simplifies and speeds file access by driving file I/O through memory rather than `read()` and `write()` system calls
- n Also allows several processes to map the same file allowing the pages in memory to be shared
- n But when does written data make it to disk?
 - | Periodically and / or at file `close()` time
 - | For example, when the pager scans for dirty pages







Memory-Mapped File Technique for all I/O

- n Some OSes use memory mapped files for standard I/O
- n Process can explicitly request memory mapping a file via mmap() system call
 - | Now file mapped into process address space
- n For standard I/O (`open()`, `read()`, `write()`, `close()`), mmap anyway
 - | But map file into kernel address space
 - | Process still does `read()` and `write()`
 - ▶ Copies data to and from kernel space and user space
 - | Uses efficient memory management subsystem
 - ▶ Avoids needing separate subsystem
- n COW can be used for read/write non-shared pages
- n Memory mapped files can be used for shared memory (although again via separate system calls)







Memory Mapped Files





Memory-Mapped Shared Memory in Windows






Allocating Kernel Memory

- n Treated differently from user memory
- n Often allocated from a free-memory pool
 - | Kernel requests memory for structures of varying sizes
 - | Some kernel memory needs to be contiguous
 - ▶ I.e. for device I/O





Buddy System

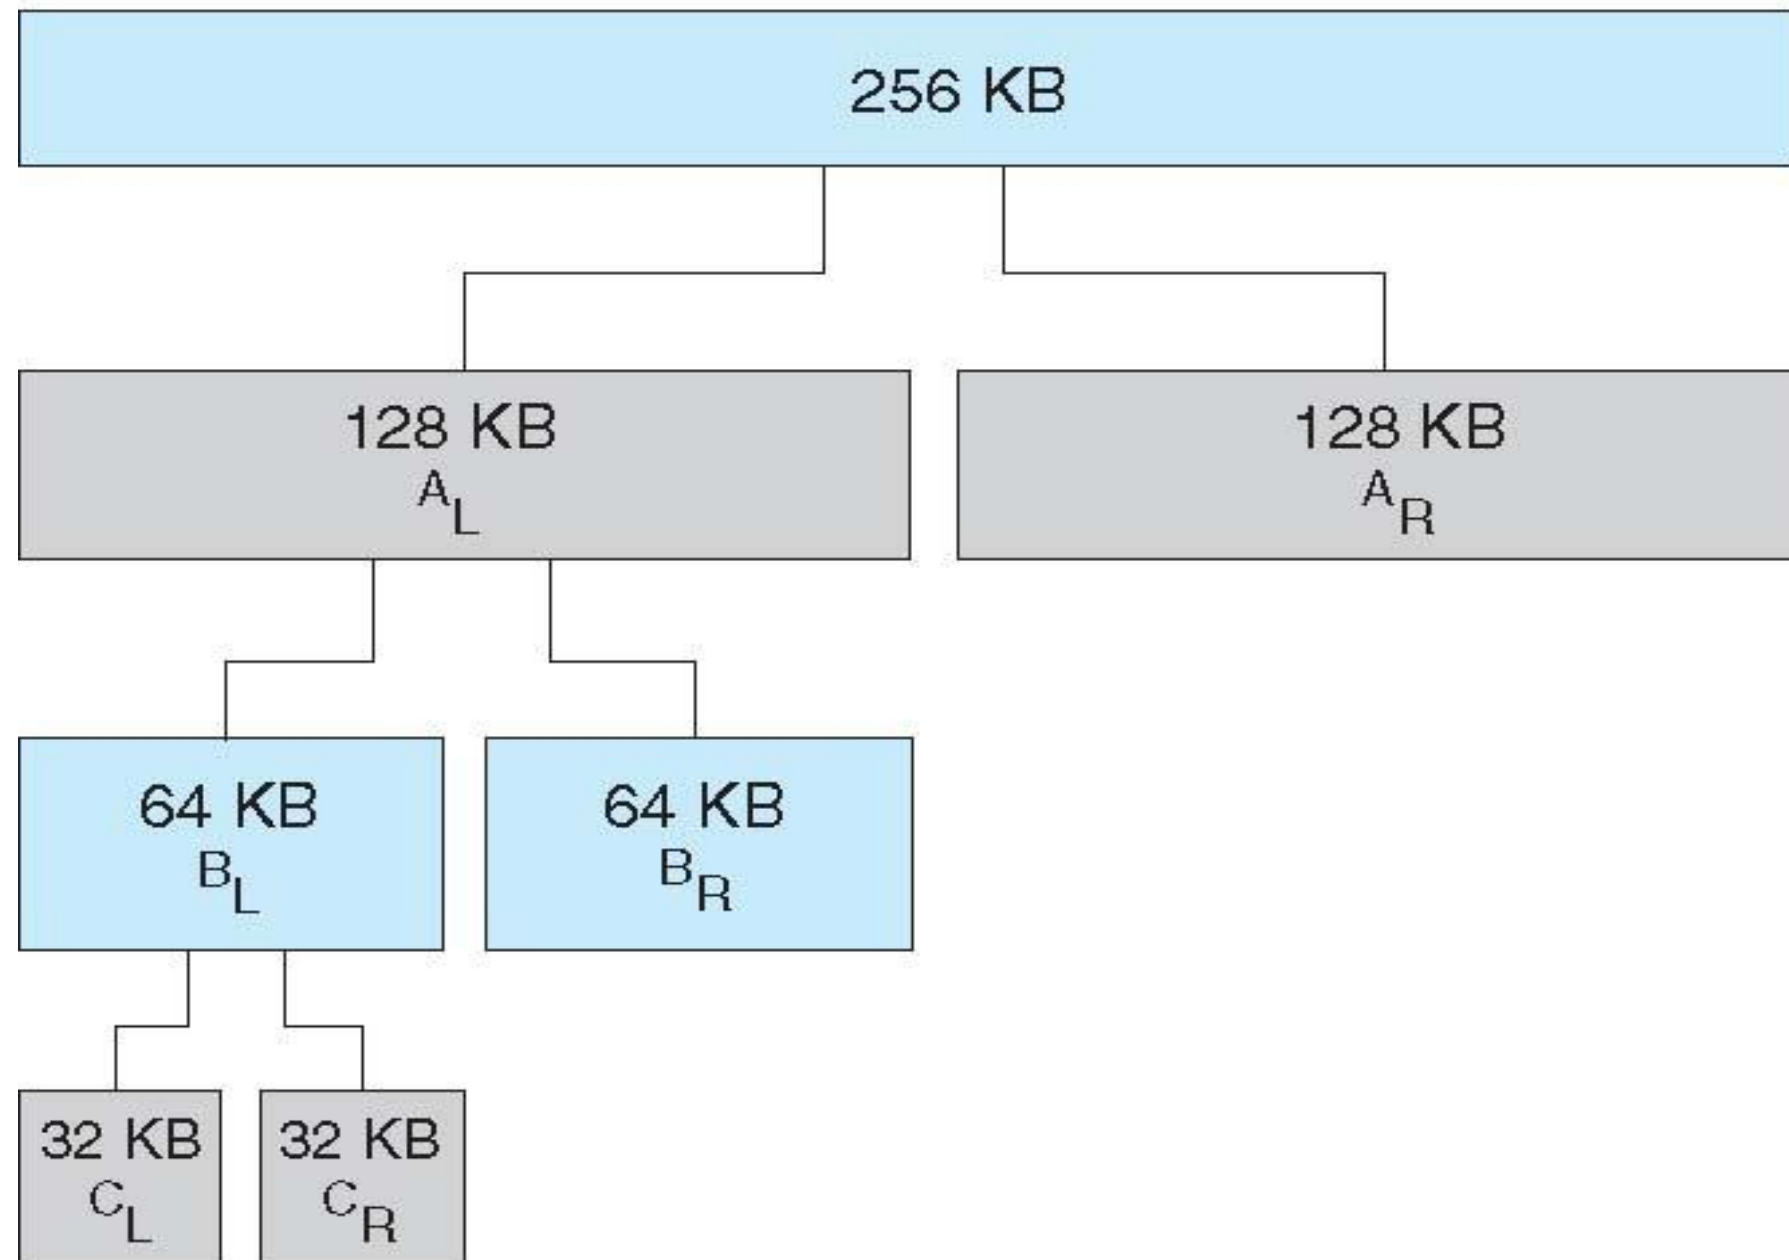
- n Allocates memory from fixed-size segment consisting of physically-contiguous pages
- n Memory allocated using power-of-2 allocator
 - | Satisfies requests in units sized as power of 2
 - | Request rounded up to next highest power of 2
 - | When smaller allocation needed than is available, current chunk split into two buddies of next-lower power of 2
 - ▶ Continue until appropriate sized chunk available
- n For example, assume 256KB chunk available, kernel requests 21KB
 - | Split into A_L and A_R of 128KB each
 - ▶ One further divided into B_L and B_R of 64KB
 - One further into C_L and C_R of 32KB each – one used to satisfy request
- n Advantage – quickly coalesce unused chunks into larger chunk
- n Disadvantage - fragmentation






Buddy System Allocator

physically contiguous pages







Slab Allocator

- n Alternate strategy
- n **Slab** is one or more physically contiguous pages
- n **Cache** consists of one or more slabs
- n Single cache for each unique kernel data structure
 - | Each cache filled with **objects** – instantiations of the data structure
- n When cache created, filled with objects marked as free
- n When structures stored, objects marked as used
- n If slab is full of used objects, next object allocated from empty slab
 - | If no empty slabs, new slab allocated
- n Benefits include no fragmentation, fast memory request satisfaction





Slab Allocation






Other Considerations -- Prepaging

- n Prepaging
 - | To reduce the large number of page faults that occurs at process startup
 - | Prepage all or some of the pages a process will need, before they are referenced
 - | But if prepaged pages are unused, I/O and memory was wasted
 - | Assume s pages are prepaged and α of the pages is used
 - ▶ Is cost of $s * \alpha$ save pages faults $>$ or $<$ than the cost of prepaging $s * (1 - \alpha)$ unnecessary pages?
 - ▶ α near zero \Rightarrow prepaging loses






Other Issues – Page Size

- n Sometimes OS designers have a choice
 - | Especially if running on custom-built CPU
- n Page size selection must take into consideration:
 - | Fragmentation
 - | Page table size
 - | Resolution
 - | I/O overhead
 - | Number of page faults
 - | Locality
 - | TLB size and effectiveness
- n Always power of 2, usually in the range 2^{12} (4,096 bytes) to 2^{22} (4,194,304 bytes)
- n On average, growing over time





Other Issues – TLB Reach

- n TLB Reach - The amount of memory accessible from the TLB
- n $\text{TLB Reach} = (\text{TLB Size}) \times (\text{Page Size})$
- n Ideally, the working set of each process is stored in the TLB
 - | Otherwise there is a high degree of page faults
- n Increase the Page Size
 - | This may lead to an increase in fragmentation as not all applications require a large page size
- n Provide Multiple Page Sizes
 - | This allows applications that require larger page sizes the opportunity to use them without an increase in fragmentation





Other Issues – Program Structure

n Program structure

- | `Int[128,128] data;`
- | Each row is stored in one page
- | Program 1

```
for (j = 0; j < 128; j++)  
    for (i = 0; i < 128; i++)  
        data[i,j] = 0;
```


128 x 128 = 16,384 page faults

| Program 2

```
for (i = 0; i < 128; i++)  
    for (j = 0; j < 128; j++)  
        data[i,j] = 0;
```

128 page faults






Other Issues – I/O interlock


- n **I/O Interlock** – Pages must sometimes be locked into memory
- n Consider I/O - Pages that are used for copying a file from a device must be locked from being selected for eviction by a page replacement algorithm





Reason Why Frames Used For I/O Must Be In Memory





Operating System Examples

- n Windows XP
- n Solaris





Windows XP

- n Uses demand paging with **clustering**. Clustering brings in pages surrounding the faulting page
- n Processes are assigned **working set minimum** and **working set maximum**
- n Working set minimum is the minimum number of pages the process is guaranteed to have in memory
- n A process may be assigned as many pages up to its working set maximum
- n When the amount of free memory in the system falls below a threshold, **automatic working set trimming** is performed to restore the amount of free memory
- n Working set trimming removes pages from processes that have pages in excess of their working set minimum






Solaris

- n Maintains a list of free pages to assign faulting processes
- n Lotsfree – threshold parameter (amount of free memory) to begin paging
- n Desfree – threshold parameter to increasing paging
- n Minfree – threshold parameter to being swapping
- n Paging is performed by pageout process
- n Pageout scans pages using modified clock algorithm
- n Scanrate is the rate at which pages are scanned. This ranges from slowscan to fastscan
- n Pageout is called more frequently depending upon the amount of free memory available
- n Priority paging gives priority to process code pages





Solaris 2 Page Scanner

