

Operating System (CSC 3150)

Tutorial 11

KAI SHEN

SCHOOL OF SCIENCE AND ENGINEERING

E-MAIL: 118010254@LINK.CUHK.EDU.CN

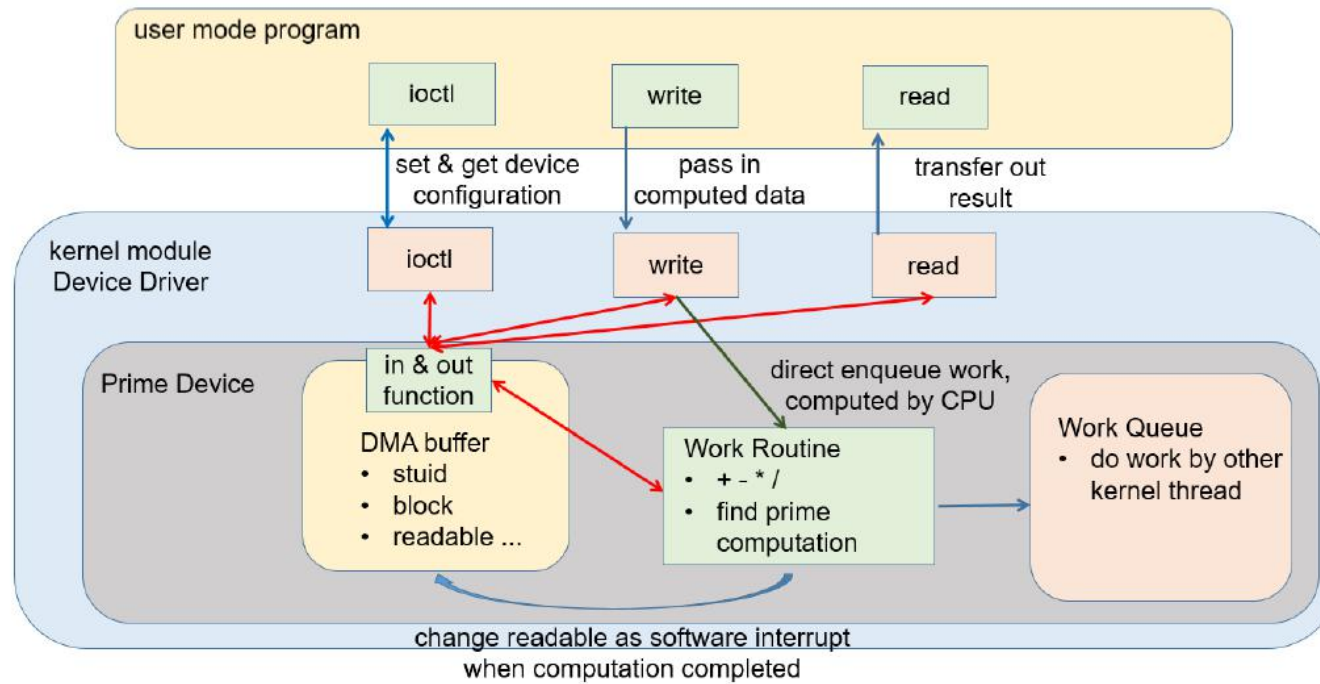
Target

In this tutorial, we will discuss Assignment 5 related functions.

- Assignment 5 Structure
- DMA buffer
- ioctl
- arithmetic
- write (blocking / non-blocking)
- read (readable)
- work
- Assignment 5 makefile and scripts
- Bonus Hints
- Assignment 5 Problems

Assignment 5 Structure

Global View:



DMA buffer

- The data is stored in DMA buffer.
- Ports being defined:
 - #define DMA_BUFSIZE 64
 - #define DMASTUIDADDR 0x0 // Student ID
 - #define DMARWOKADDR 0x4 // RW function complete
 - #define DMAIOCOKADDR 0x8 // ioctl function complete
 - #define DMAIRQOKADDR 0xc // ISR function complete
 - #define DMACOUNTADDR 0x10 // interrupt count function complete
 - #define DMAANSADDR 0x14 // Computation answer
 - #define DMAREADABLEADDR 0x18 // READABLE variable for synchronize
 - #define DMABLOCKADDR 0x1c // Blocking or non-blocking IO
 - #define DMAOPCODEADDR 0x20 // data.a opcode
 - #define DMAOPERANDBADDR 0x21 // data.b operand1
 - #define DMAOPERANDCADDR 0x25 // data.c operand2

DMA buffer

- When doing data transfer within kernel, you could use in/out function. “c/s/i” depends on what type of data you want to read or write.
 - In and out function to write/read into/from DMA buffer. (Already defined in template)
 - void **myoutc**(unsigned char data,unsigned short int port)
{ *(volatile unsigned char*)(dma_buf+**port**) = data; }
 - void **myouts**(unsigned short data,unsigned short int port)
{ *(volatile unsigned short*)(dma_buf+port) = data;}
 - void **myouti**(unsigned int data,unsigned short int port)
{ *(volatile unsigned int*)(dma_buf+port) = data; }
 - unsigned char **myinc**(unsigned short int port)
{ return *(volatile unsigned char*)(dma_buf+port); }
 - unsigned short **myins**(unsigned short int port)
{ return *(volatile unsigned short*)(dma_buf+port); }
 - unsigned int **myini**(unsigned short int port)
{ return *(volatile unsigned int*)(dma_buf+port);}
- Write data into DMA buffer with specific port.
- Read data from DMA buffer with specific port.

DMA buffer

- Demo usage of in and out function:
 - In user program, use ioctl to set I/O mode:
`int ret = 0;`
`ioctl(fd, HW5_IOCSETBLOCK, &ret);`
 - Transfer data from user to kernel:
`int value;`
`get_user(value, (int *)arg);`
 - Store I/O mode to DMA buffer in kernel:
`myouti(value, DMABLOCKADDR);`
 - Get I/O mode from DMA buffer in kernel:
`int IOMode = myini(DMABLOCKADDR);`

ioctl

- Set and get device configuration
- Masked labels: (defined in “ioc_hw5.h”)
 - (HW5_IOC_SETSTUID) Set student ID: printk your student ID
 - (HW5_IOCSETRWOK) Set if RW OK: printk OK if you complete R/W function
 - (HW5_IOCSETIOCOK) Set if ioctl OK: printk OK if you complete ioctl function
 - (HW5_IOCSETIRQOK) Set if IRQ OK: printk OK if you complete bonus
 - (HW5_IOCSETBLOCK) Set blocking or non-blocking: set write function mode
 - (HW5_IOCWAITREADABLE) Wait if readable now (synchronize function): used before read to confirm it can read answer now when use non-blocking write mode.

ioctl

- In user program, when ioctl is called, it will map to drv_ioctl in kernel. (If you've defined operation ioctl as drv_ioctl when adding the cdev.)

```
// cdev file_operations
static struct file_operations fops = {
    owner: THIS_MODULE,
    read: drv_read,
    write: drv_write,
    unlocked_ioctl: drv_ioctl,
    open: drv_open,
    release: drv_release,
};
```

- In kernel, when you received different command from user program, you need to use the in and out function to change the configurations in DMA buffer.

arithmetic

- In user program, use arithmetic function to trigger read and write.
- When write is triggered, it will transfer data into kernel. At the same time, it will put the arithmetic work into work queue.

arithmetic

- For blocking write, the work queue will be forced to wait the termination of computation. So we could read the result when write completed.
- For non-blocking write, the work routine will be continued. So in user program, we use `ioctl` to check device's readable configuration before read the result. This is for synchronization and ensure you've read the correct answer.

```
/******Blocking IO******/
printf("Blocking IO\n");
ret = 1;
if (ioctl(fd, HW5_IOCSETBLOCK, &ret) < 0) {
    printf("set blocking failed\n");
    return -1;
}

write(fd, &data, sizeof(data));

printf("testing\n");

//Do not need to synchronize
//But need to wait computation completed

read(fd, &ret, sizeof(int));

printf("ans=%d ret=%d\n\n", ans, ret);
/*******/
```

```
/******Non-Blocking IO******/
printf("Non-Blocking IO\n");
ret = 0;
if (ioctl(fd, HW5_IOCSETBLOCK, &ret) < 0) {
    printf("set non-blocking failed\n");
    return -1;
}

printf("Queueing work\n");
write(fd, &data, sizeof(data));

//Can do something here
//But cannot confirm computation completed
printf("testing\n");

printf("Waiting\n");
//synchronize function
ioctl(fd, HW5_IOCWAITREADABLE, &readable);

if(readable==1){
    printf("Can read now.\n");
    read(fd, &ret, sizeof(int));
}
printf("ans=%d ret=%d\n\n", ans, ret);
/*******/
```

write (blocking / non-blocking)

- In user program, write the data into your device.
- In kernel, transfer data into DMA buffer.
- Check the device I/O mode.
- Use “INIT_Work” to define what function (drv_arithmetic) should be executed in work routine.
- Basing on I/O mode (blocking or non-blocking), place the work into work queue.

read (readable)

- In user program, use `ioctl` to check readable flag to synchronize the result when there is non-blocking write.
- In user program, read the computed result from device.
- In kernel, read the result from DMA buffer.
- Clean the result and set readable as false.
- Readable setting is checked via `drv_ioctl`. If it is non-readable, then just sleep and waiting.

drv_arithmetic

- In kernel, drv_arithmetic is to execute the computation when the work is in queue.
- It should complete '+', '-', '*', '/' and 'p' computations basing on the data being stored in DMA buffer.
- It should check blocking and non-blocking setting from DMA buffer.
- If the device is executing non-blocking write, set the readable as true when completed computation.

work

- Define a work: `struct work_struct *work`
- `INIT_WORK(work, func)`: initialize work. Func is defined what function need to be executed for this work.
- `schedule_work(work)`: put work task in global workqueue.
- `flush_scheduled_work()`: It is to flush the kernel-global work queue. Forces execution of the kernel-global workqueue and blocks until its completion.
- Use work to control block and non-blocking write.

work example (kernel)

```
91 static int __init init_modules(void)
92 {
93     dev_t dev;
94
95     printk("%s:()%s():.....Start.....\n", PREFIX_TITLE, __func__);
96
97     dev_cdev = cdev_alloc();
98
99     // Register chrdev
100     if(alloc_chrdev_region(&dev, DEV_BASEMINOR, DEV_COUNT, DEV_NAME) < 0) {
101         printk(KERN_ALERT "Register chrdev failed!\n");
102         return -1;
103     } else {
104         printk("%s:()%s(): register chrdev(%i,%i)\n", PREFIX_TITLE, __func__, MAJOR(dev), MINOR(dev));
105     }
106
107     dev_major = MAJOR(dev);
108     dev_minor = MINOR(dev);
109
110     // Init cdev
111     dev_cdev->ops = &fops;
112     dev_cdev->owner = THIS_MODULE;
113
114     if(cdev_add(dev_cdev, dev, 1) < 0) {
115         printk(KERN_ALERT "Add cdev failed!\n");
116         return -1;
117     }
118
119
120     // Alloc work routine
121     work = kzalloc(sizeof(typeof(*work)), GFP_KERNEL);
122
123     return 0;
124 }
125
126 static void __exit exit_modules(void) {
127
128     // Delete char device
129     unregister_chrdev_region(MKDEV(dev_major, dev_minor), DEV_COUNT);
130     cdev_del(dev_cdev);
131
132     // Free work routine
133     kfree(work);
134     printk("%s:()%s(): unregister chrdev\n", PREFIX_TITLE, __func__);
135     printk("%s:()%s():.....End.....\n", PREFIX_TITLE, __func__);
136 }
137 }
```

work example (kernel)

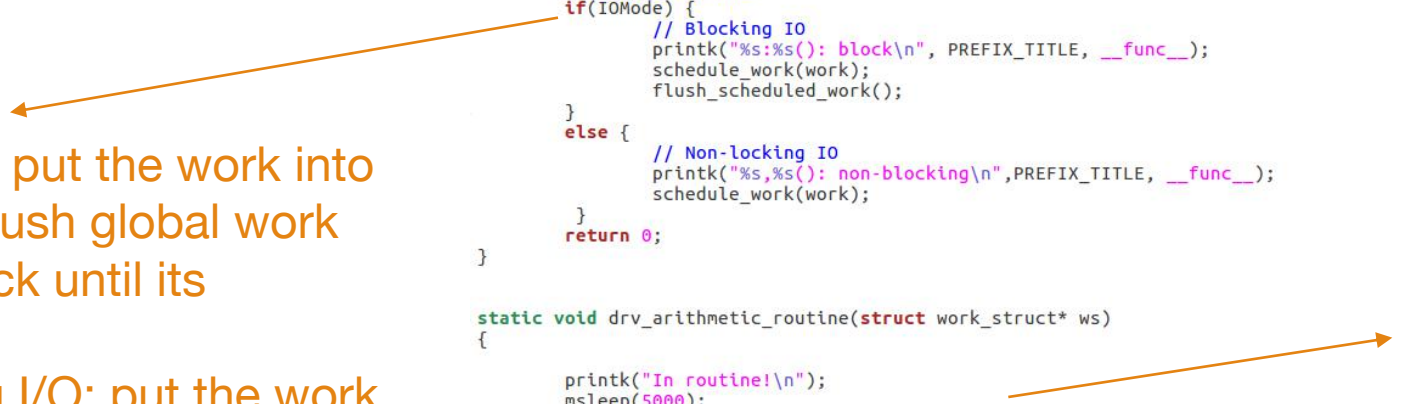
```
29 // File Operations
30 static int drv_open(struct inode*, struct file*);
31 static ssize_t drv_write(struct file *filp, const char __user *buffer, size_t, loff_t*);
32 static int drv_release(struct inode*, struct file*);
33 static struct file_operations fops = {
34     owner: THIS_MODULE,
35     write: drv_write,
36     open: drv_open,
37     release: drv_release,
38 };

static ssize_t drv_write(struct file *filp, const char __user *buffer, size_t ss, loff_t* lo) {
    int IOMode;
    get_user(IOMode, (int *) buffer);
    printk("%s:%s(): IO Mode is %d\n", PREFIX_TITLE, __func__, IOMode);

    INIT_WORK(work, drv_arithmetic_routine);

    // Decide io mode
    if(IOMode) {
        // Blocking IO
        printk("%s:%s(): block\n", PREFIX_TITLE, __func__);
        schedule_work(work);
        flush_scheduled_work();
    }
    else {
        // Non-locking IO
        printk("%s:%s(): non-blocking\n", PREFIX_TITLE, __func__);
        schedule_work(work);
    }
    return 0;
}

static void drv_arithmetic_routine(struct work_struct* ws)
{
    printk("In routine!\n");
    msleep(5000);
    printk("Work routine completed!\n");
}
```



Blocking I/O: put the work into queue, and flush global work queue to block until its completion.

Non blocking I/O: put the work into queue.

For this work, let it sleep for 5 seconds.

work example (user program)

```
/******Blocking IO*****/
printf("Blocking I/O\n");
ret = 1;
write(fd, &ret, sizeof(ret));

//Do not need to synchronize

printf("Blocking I/O completed!\n");
/*******/

/******Non-Blocking IO*****/
printf("Non-Blocking I/O\n");
ret = 0;
write(fd, &ret, sizeof(ret));


//Can do something here
//But cannot confirm computation completed

printf("Non-Blocking I/O is still executing in kernel!\n");
/*******/
```

work example (output)

User program output

```
[11/20/19]seed@VM:~/.../Work$ ./test
.....Start.....
Blocking I/O
```



For blocking write, you will see it will not continue print message until the work is executed completely.

Kernel log output (display immediately when user program terminates)

```
[ 2956.005097] Tutorial_11:init_modules():.....Start.....
[ 2956.005099] Tutorial_11:init_modules(): register chrdev(245,0)
[ 2960.324039] Tutorial_11:drv_open(): device open
[ 2960.324042] Tutorial_11:drv_write(): IO Mode is 1
[ 2960.324043] Tutorial_11:drv_write(): block
[ 2960.324046] In routine!
[ 2965.331478] Work routine completed!
[ 2965.331533] Tutorial_11:drv_write(): IO Mode is 0
[ 2965.331534] Tutorial_11:drv_write(): non-blocking
[ 2965.331897] Tutorial_11:drv_release(): device close
[ 2965.336558] In routine!
```

```
[11/20/19]seed@VM:~/.../Work$ ./test
.....Start.....
Blocking I/O
Blocking I/O completed!
Non-Blocking I/O
Work routine is executed in progress within kernel!
.....End.....
[11/20/19]seed@VM:~/.../Work$
```

Kernel log output (display 5 seconds later)

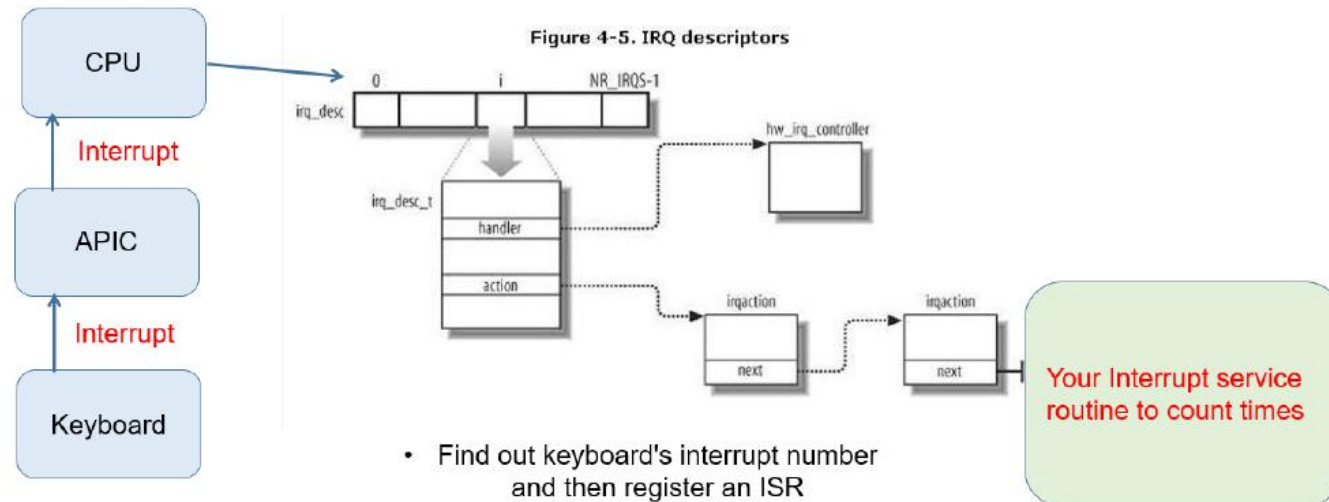
```
[2956.005097] Tutorial_11:init_modules():.....Start.....
[2956.005099] Tutorial_11:init_modules(): register chrdev(245,0)
[2960.324039] Tutorial_11:drv_open(): device open
[2960.324042] Tutorial_11:drv_write(): IO Mode is 1
[2960.324043] Tutorial_11:drv_write(): block
[2960.324046] In routine!
[2965.331478] Work routine completed!
[2965.331533] Tutorial_11:drv_write(): IO Mode is 0
[2965.331534] Tutorial_11:drv_write(): non-blocking
[2965.331897] Tutorial_11:drv_release(): device close
[2965.336558] In routine!
[2970.448867] Work routine completed!
```

Assignment 5 makefile and scripts

- Command “**make**” includes executions:
 - Build kernel module “mydev.ko”
 - Insert kernel module “sudo insmod mydev.ko”
- For first time active your device:
 - Check available MAJOR and MINOR number. (When initialize your kernel module, use `alloc_chrdev_region()` and `MAJOR()` / `MINOR()` to check. After insert your module, type “**dmesg**” to check the result.)
 - create a file node for “mydev”. Run “**sudo ./mkdev.sh MAJOR MINOR**”.
 - You don’t need to create this file node every time. Just create it if it does not exist. (You can run “**sudo ./rmdev.sh**” to remove this file node.)
- Run test “**./test**”:
- Command “**make clean**” includes executions:
 - Remove kernel module “sudo rmmod mydev.ko”
 - Clean test executable file “rm test”
 - Display kernel message including “OS_AS5” (“**dmesg | grep OS_AS5**”)

Bonus Hints

Global View (Bonus)



Bonus Hints

- `request_irq (unsigned int irq,
irq_handler_t handler,
unsigned long irqflags,
const char * devname,
void * dev_id);`
 - `irq`: Interrupt line to allocate (Hints: Can check from “watch -n 1 cat /proc/interrupts”)
 - `handler`: Function to be called when the IRQ occurs. (Hints: count interrupt times when IRQ occurs)
 - `irqflags`: Interrupt type flags (Hints: `IRQF_SHARED` Interrupt is shared)
 - `devname`: An ascii name for the claiming device. (Hints: Once your kernel module is inserted/removed, it will auto be displayed under “watch -n 1 cat /proc/interrupts”. Define any name you want to show.)
 - `dev_id`: A cookie passed back to the handler function.

Bonus Hints

- `free_irq (unsigned int irq,
 void * dev_id);`
 - `irq`: Interrupt line to free
 - `dev_id`: Device identity to free.

- `irq_handler_t` handler

- <https://elixir.bootlin.com/linux/v4.10.14/source/include/linux/interrupt>
- `typedef irqreturn_t (*irq_handler_t)(int, void *);`
- <https://elixir.bootlin.com/linux/v4.10.14/source/include/linux/irqreturn.h>

```
/**
 * enum irqreturn
 * @IRQ_NONE      interrupt was not from this device or was not handled
 * @IRQ_HANDLED   interrupt was handled by this device
 * @IRQ_WAKE_THREAD handler requests to wake the handler thread
 */
enum irqreturn {
    IRQ_NONE      = (0 << 0),
    IRQ_HANDLED   = (1 << 0),
    IRQ_WAKE_THREAD = (1 << 1),
};
```

- More detailed usage for irq handler
 - Understanding Linux Kernel, Page 591 - 595

Bonus Hints

- Run command “watch -n 1 cat /proc/interrupts”
 - IRQ_NUM should be defined as 1 for below case. (Share interrupt with i8042, the keyboard interrupt)

```
Terminal
Every 1.0s: cat /proc/interrupts      Tue Dec  4 01:49:36 2018

    CPU0
 0:   31      XT-PIC  timer
 1: 2226      XT-PIC  i8042
 2:    0      XT-PIC  cascade
 8:    0      XT-PIC  rtc0
 9: 8263      XT-PIC  acpi, enp0s3
10:  29      XT-PIC  ohci hcd:usb1, vboxvideo
11: 75693     XT-PIC  ahci[0000:00:0d.0], vboxguest, snd_intel8x0
12:  7488     XT-PIC  i8042
14:    0      XT-PIC  ata_piix
15: 15906     XT-PIC  ata_piix
NMI:    0      Non-maskable interrupts
LOC: 1985609   Local timer interrupts
SPU:    0      Spurious interrupts
PMI:    0      Performance monitoring interrupts
IWI:    0      IRQ work interrupts
RTR:    0      APIC ICR read retries
RES:    0      Rescheduling interrupts
CAL:    0      Function call interrupts
TLB:    0      TLB shootdowns
TRM:    0      Thermal event interrupts
THR:    0      Threshold APIC interrupts
```

Bonus Hints

- Inserted mydev.ko

```
Terminal
Every 1.0s: cat /proc/interrupts      Tue Dec  4 01:50:13 2018

          CPU0
0:         31   XT-PIC  timer
1:        2236   XT-PIC  i8042, myinterrupts
```

- After run test

```
Terminal
Every 1.0s: cat /proc/interrupts      Tue Dec  4 01:51:01 2018

          CPU0
0:         31   XT-PIC  timer
1:        2250   XT-PIC  i8042, myinterrupts
```

- Remove mydev.ko

```
Terminal
Every 1.0s: cat /proc/interrupts      Tue Dec  4 01:51:32 2018

          CPU0
0:         31   XT-PIC  timer
1:        2272   XT-PIC  i8042
```

- Interrupt counting

```
[16052.017861] OS_AS5:drv_ioctl(): wait readable 1
[16052.017869] OS_AS5:drv_read(): ans = 225077
[16052.017987] OS_AS5:drv_release(): device close
[16086.218416] OS_AS5:exit_modules(): interrupt count=36
[16086.218417] OS_AS5:exit_modules(): free dma buffer
```


Assignment 5 Problems

- “Command not found” for “sudo ./mkdev.sh Major Minor”.
 - Do not use the script anymore, type the command directly.
 - `sudo mknod /dev/mydev c major minor`
 - `sudo chmod 666 /dev/mydev`
 - `ls -l /dev/mydev`

Assignment 5 Problems

- `get_user` / `put_user` cannot transfer data with type `DataIn`
 - This macro copies a single simple variable from user space to kernel space.
 - It supports simple types like `char` and `int`, but not larger data types like structures or arrays.
 - Some students may get killed error when running test program if transferring `DataIn` via `get_user` / `put_user`.
 - For data with type `DataIn`, you can transfer via buffer, second parameter of the `drv_write` / `drv_read`.
 - Or you could use `copy_from_user` / `copy_to_user`.

Assignment 5 Problems

- How to change device readable setting?
 - When the device is unreadable, use a while loop with msleep in kernel.
 - Use the while loop to keep checking readable setting.
 - If unreadable, use msleep to let this work to be slept.
 - In this case, the CPU is available to continue computation.
 - Once the computation completes, change readable setting.
 - Then in while loop, if it detects the readable setting is updated, it will stop looping.

Assignment 5 Problems

- When will `drv_release` be called?
 - If `close()` is called in user program, `drv_release()` will be called in kernel.
 - If `close()` is not called in user program, when that process is terminated (`./test` completes execution), `close()` will be automatically called.

Assignment 5 Problems

- Bonus: interrupt counting will not always be 36.
 - It depend on what you've typed in via keyboard.
 - You only need to calculate the interrupt numbers during mydev.ko being inserted.
 - You could verify the number via “watch -n 1 cat /proc/interrupts”.
 - Sometimes, even if you only press once in keyboard, the interrupt will occur twice or multiple times.

Reference

- Work queues
 - <https://www.ibm.com/developerworks/library/l-tasklets/index.html>
 - <https://linuxtv.org/downloads/v4l-dvb-internals/device-drivers/ch01s06.html>

- Get user / Put user
 - <https://www.fsl.cs.sunysb.edu/kernel-api/re244.html>
 - <https://www.fsl.cs.sunysb.edu/kernel-api/re245.html>

- Request and free irq
 - <https://elixir.bootlin.com/linux/v4.10.14/source/include/linux/interrupt.h>
 - <https://www.fsl.cs.sunysb.edu/kernel-api/re667.html>
 - <https://www.kernel.org/doc/html/docs/kernel-api/API-free-irq.html>

Thank you

